

## Outline

In this notebook I am going to describe the major steps toward creating a CNN from scratch as well as using transfer learning to complete the Dog Breed classification project.

### 1. Initialization ¶

### 2. Define a basic CNN to start

### 3. Data Exploration and need for Augmentation

### 4. Tune CNN and hyper parameters

### 5. Transfer learning

### 6. Conclusion

## Initialization:

The datasets are provided for us and divided into test, train and validation each have 113 folders corresponding to 113 dog breeds. The label distribution seems acceptable (mean 6.272727 std 1.712509 min 3.000000 max 10.000000), although it could be better considering number of samples.

<https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/dogImages.zip> (<https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/dogImages.zip>)

<https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/lfw.zip> (<https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/lfw.zip>)

```
In [69]: import numpy as np
from glob import glob

# load filenames for human and dog images
human_files = np.array(glob("lfw/**/*.jpg"))
dog_files = np.array(glob("dogImages/**/*.jpg"))

# print number of images in each dataset
print('There are %d total human images.' % len(human_files))
print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.  
There are 8351 total dog images.

```
In [70]: dog_files_test = np.array(glob("dogImages/test/**/*.jpg"))
dog_files_train = np.array(glob("dogImages/train/**/*.jpg"))
dog_files_valid = np.array(glob("dogImages/valid/**/*.jpg"))

# print number of images in each dataset
print('There are %d total dog images in test.' % len(dog_files_test))
print('There are %d total dog images in train.' % len(dog_files_train))
print('There are %d total dog images in valid.' % len(dog_files_valid))
```

There are 836 total dog images in test.  
There are 6680 total dog images in train.  
There are 835 total dog images in valid.

```
In [71]: import torch
from PIL import Image
import torchvision.transforms as transforms
import json
from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True
```

```
In [87]: import torchvision.models as models
import torch
import torchvision.models as models
# define VGG16 model
VGG16 = models.vgg16(pretrained=True)
# check if CUDA is available
use_cuda = torch.cuda.is_available()
if not use_cuda:
    device = "cpu"
    print('Training on CPU')
else:
    device = torch.device("cuda:0")
    print('Training on GPU ...')
    print("Using", torch.cuda.get_device_name(device))
# move model to GPU if CUDA is available
# if use_cuda:
    # VGG16 = VGG16.cuda()
VGG16 = VGG16.to(device)
```

```
Training on GPU ...
Using Tesla V100-SXM2-16GB
```

## Define a CNN for classification

\*Network architecture

There are 3 sets of two convolutional layer followed by a max pooling to control computational complexity, the initial depth is 32 which would be doubled after each set. It is inspired by VGG network architecture. dropout is applied (to avoid over fitting) before send the result to a 1000 flat output layer (correspond to dog breeds). I am going to use this architecture with arbitrary learning rate of 0.01 to see how it performs on a non-augmented data, later if it is needed I am going to augment and compare the results. after that I am going to tune the model itself.

<https://arxiv.org/abs/1409.1556> (<https://arxiv.org/abs/1409.1556>)

```

In [88]: import torch.nn as nn
import torch.nn.functional as F

# define the CNN architecture
class Net(nn.Module):
    def __init__(self, n_classes, depth_1 = 32):
        super(Net, self).__init__()
        # Keep track of things
        depth_2 = depth_1 * 2
        depth_3 = depth_2 * 2
        # Max pooling layer
        self.pool = nn.MaxPool2d(2,2)
        # Conv set 1
        self.conv1_1 = nn.Conv2d(3,depth_1,3,stride = 1,padding = 1)
        self.conv1_2 = nn.Conv2d(depth_1,depth_1,3,stride = 1,padding
= 1)

        self.bn1_1 = nn.BatchNorm2d(depth_1)
        self.bn1_2 = nn.BatchNorm2d(depth_1)
        # Conv set 2
        self.conv2_1 = nn.Conv2d(depth_1,depth_2,3,stride = 1,padding
= 1)

        self.conv2_2 = nn.Conv2d(depth_2,depth_2,3,stride = 1,padding
= 1)

        self.bn2_1 = nn.BatchNorm2d(depth_2)
        self.bn2_2 = nn.BatchNorm2d(depth_2)
        # Conv set 3
        self.conv3_1 = nn.Conv2d(depth_2,depth_3,3,stride = 1,padding
= 1)

        self.conv3_2 = nn.Conv2d(depth_3,depth_3,3,stride = 1,padding
= 1)

        self.bn3_1 = nn.BatchNorm2d(depth_3)
        self.bn3_2 = nn.BatchNorm2d(depth_3)
        # Output correspond to number of dogs breeds
        self.fc_out = nn.Linear(depth_3,n_classes)
        # Initialize weights
        nn.init.kaiming_normal_(self.conv1_1.weight, nonlinearity='re
lu')
        nn.init.kaiming_normal_(self.conv1_2.weight, nonlinearity='re
lu')
        nn.init.kaiming_normal_(self.conv2_1.weight, nonlinearity='re
lu')
        nn.init.kaiming_normal_(self.conv2_2.weight, nonlinearity='re
lu')
        nn.init.kaiming_normal_(self.conv3_1.weight, nonlinearity='re
lu')
        nn.init.kaiming_normal_(self.conv3_2.weight, nonlinearity='re
lu')

    def forward(self, x):
        # Conv 1
        x = F.relu(self.bn1_1(self.conv1_1(x)))
        x = F.relu(self.bn1_2(self.conv1_2(x)))
        x = self.pool(x)
        # Conv 2
        x = F.relu(self.bn2_1(self.conv2_1(x)))
        x = F.relu(self.bn2_2(self.conv2_2(x)))

```

```
x = self.pool(x)
# Conv 3
x = F.relu(self.bn3_1(self.conv3_1(x)))
x = F.relu(self.bn3_2(self.conv3_2(x)))
x = self.pool(x)
# reduce dimention
x = x.view(x.size(0),x.size(1),-1)
# And now max global pooling
x = x.max(2)[0]
# Output
x = self.fc_out(x)
return x
# instantiate the CNN
model_scratch = Net(n_classes)

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()
```

## Data Exploration and need for Augmentation

There are images with different size that we need to crop/resize to the same size at first step. As later I am going to compare it with pretrained model, I am going to apply the same size and normalization that they are trained with here as well. Now before going further, I want to observe how it is the learning/validation error ratio is for 10 epoches. I am going to augment and check it again to see if there is a meaningful improvement on overfitting issue. specially given the fact that we don't have much of a training data set.

```

In [74]: import os
from torchvision import datasets

### No Augmentation Performance review
transform_resize = 224
transform_crop = 224
data_directory = "dogImages"

print("load image data ... ")
# define transforms for the training data and testing data
train_transforms = transforms.Compose([transforms.CenterCrop(transform_crop), transforms.ToTensor()])

test_transforms = transforms.Compose([transforms.CenterCrop(transform_crop), transforms.ToTensor()])

# pass transforms in here, then run the next cell to see how the transforms look
train_data = datasets.ImageFolder( data_directory + '/train', transform=train_transforms )
test_data = datasets.ImageFolder( data_directory + '/test', transform=test_transforms )
valid_data = datasets.ImageFolder( data_directory + '/valid', transform=test_transforms )

# ---- print out some data stats ----
print(' Number of train images: ', len(train_data))
print(' Number of test images: ', len(test_data))
print(' Number of valid images: ', len(valid_data))
# -----

trainloader = torch.utils.data.DataLoader( train_data, batch_size=32, shuffle=True )
testloader = torch.utils.data.DataLoader( test_data, batch_size=16 )
validloader = torch.utils.data.DataLoader( valid_data, batch_size=16 )

# create dictionary for all loaders in one
loaders_scratch = {}
loaders_scratch['train'] = trainloader
loaders_scratch['valid'] = validloader
loaders_scratch['test'] = testloader
n_classes = len(train_data.classes)

load image data ...
Number of train images: 6680
Number of test images: 836
Number of valid images: 835

```

```
In [89]: import torch.optim as optim
from torch.optim.lr_scheduler import ReduceLRonPlateau

### TODO: select loss function
# criterion_scratch = None

param_learning_rate = 0.01 # I tweaked it manually
criterion_scratch = nn.CrossEntropyLoss() # Based on pytorch recommendation for multi class classification with high precision and recall

### TODO: select optimizer
optimizer_scratch = optim.SGD(model_scratch.parameters(), lr=param_learning_rate)
```

```

In [76]: # the following import is required for training to be robust to truncated images
from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True

def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf

    for epoch in range(1, n_epochs+1):
        # initialize variables to monitor training and validation loss
        train_loss = 0.0
        valid_loss = 0.0

        #####
        # train the model #
        #####
        model.train()
        for batch_idx, (data, target) in enumerate(loaders['train']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            ## find the loss and update the model parameters accordingly

            ## record the average training loss, using something like
            ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
            optimizer.zero_grad()
            # forward pass: compute predicted outputs by passing inputs to the model
            output = model(data)
            # calculate the batch loss
            loss = criterion(output, target)
            # backward pass: compute gradient of the loss with respect to model parameters
            loss.backward()
            # perform a single optimization step (parameter update)
            optimizer.step()
            # update training loss
            #train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
            train_loss += loss.item()*data.size(0)
            #####
            # validate the model #
            #####
            model.eval()
            for batch_idx, (data, target) in enumerate(loaders['valid']):
                # move to GPU
                if use_cuda:
                    data, target = data.cuda(), target.cuda()
                ## update the average validation loss
                output = model(data)
                # calculate the batch loss

```



```

        loss = criterion(output, target)
        # update average validation loss
        valid_loss += loss.item() * data.size(0)

    # print training/validation statistics
    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
        epoch,
        train_loss,
        valid_loss
    ))

    ## TODO: save the model if validation loss has decreased
    if valid_loss <= valid_loss_min:
        #print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(valid_loss_min, valid_loss))
        print(' Saving model ...')
        torch.save(model.state_dict(), save_path)
        valid_loss_min = valid_loss
    else:
        print("")

    # return trained model
    return model

# train the model
model_scratch = train(10, loaders_scratch, model_scratch, optimizer_scratch,
                      criterion_scratch, use_cuda, 'none_augment/model_noaugmentation_scratch.pt')

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('none_augment/model_noaugmentation_scratch.pt'))

```

```

Epoch: 1      Training Loss: 33439.447720      Validation Loss: 406
8.734119
  Saving model ...
Epoch: 2      Training Loss: 31758.934265      Validation Loss: 395
0.578050
  Saving model ...
Epoch: 3      Training Loss: 30885.299099      Validation Loss: 386
9.568702
  Saving model ...
Epoch: 4      Training Loss: 30106.976276      Validation Loss: 381
4.408498
  Saving model ...
Epoch: 5      Training Loss: 29459.478176      Validation Loss: 379
3.919446
  Saving model ...
Epoch: 6      Training Loss: 28843.910347      Validation Loss: 372
6.464614
  Saving model ...
Epoch: 7      Training Loss: 28305.996410      Validation Loss: 367
7.568949
  Saving model ...
Epoch: 8      Training Loss: 27765.342726      Validation Loss: 363
3.714413
  Saving model ...
Epoch: 9      Training Loss: 27162.707621      Validation Loss: 363
9.395047

Epoch: 10     Training Loss: 26672.503998      Validation Loss: 358
6.197377
  Saving model ...

```

**Out[76]:** <All keys matched successfully>

## Note1:

It appears that after epoch 8 validation loss stop decreasing and it would diverges from training loss, to avoid this I came up with an augemtaion like bellow for model to learn invarient representation of a dog, next I am going to test the laerning/validation error ratio for 10 epoches to see if there are improvment. Resize: centerCrop: random Horizontal and random vertical flip: random rotation: normilization

```

In [78]: import os
from torchvision import datasets

### TODO: Write data loaders for training, validation, and test sets
## Specify appropriate transforms, and batch_sizes
transform_resize = 224
transform_crop = 224
data_directory = "dogImages"

print("load image data ... ")
# define transforms for the training data and testing data
train_transforms = transforms.Compose([transforms.Resize(transform_re
size),
                                     transforms.CenterCrop(transform
m_crop),
                                     transforms.RandomHorizontalFli
p(),
                                     transforms.RandomVerticalFlip
(),
                                     transforms.RandomRotation(20),
                                     transforms.ToTensor(),
                                     transforms.Normalize([0.485,
0.456, 0.406],
                                                         [0.229,
0.224, 0.225])])

test_transforms = transforms.Compose([transforms.Resize(transform_res
ize),
                                     transforms.CenterCrop(transform
_crop),
                                     transforms.ToTensor(),
                                     transforms.Normalize([0.485, 0.
456, 0.406],
                                                         [0.229, 0.
224, 0.225])])

# pass transforms in here, then run the next cell to see how the tran
sforms look
train_data = datasets.ImageFolder( data_directory + '/train', transfo
rm=train_transforms )
test_data = datasets.ImageFolder( data_directory + '/test', transform
=test_transforms )
valid_data = datasets.ImageFolder( data_directory + '/valid', transfo
rm=test_transforms )

# ---- print out some data stats ----
print('  Number of train images: ', len(train_data))
print('  Number of test images: ', len(test_data))
print('  Number of valid images: ', len(valid_data))
# -----

trainloader = torch.utils.data.DataLoader( train_data, batch_size=32,
shuffle=True )
testloader = torch.utils.data.DataLoader( test_data, batch_size=16 )
validloader = torch.utils.data.DataLoader( valid_data, batch_size=16

```

```
)

# create dictionary for all loaders in one
loaders_scratch = {}
loaders_scratch['train'] = trainloader
loaders_scratch['valid'] = validloader
loaders_scratch['test'] = testloader
n_classes = len(train_data.classes)

load image data ...
    Number of train images: 6680
    Number of test images: 836
    Number of valid images: 835
```

```
In [79]: model_scratch = train(10, loaders_scratch, model_scratch, optimizer_s
        scratch,
        criterion_scratch, use_cuda, 'none_augment/model_augmentation_scratch.pt')

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('none_augment/model_augmentation_scratch.pt'))
```

```
Epoch: 1      Training Loss: 29126.578182      Validation Loss: 361
9.310997
    Saving model ...
Epoch: 2      Training Loss: 28206.389803      Validation Loss: 351
2.117984
    Saving model ...
Epoch: 3      Training Loss: 27685.631363      Validation Loss: 350
8.756071
    Saving model ...
Epoch: 4      Training Loss: 27081.482113      Validation Loss: 345
2.839887
    Saving model ...
Epoch: 5      Training Loss: 26545.053951      Validation Loss: 338
5.909094
    Saving model ...
Epoch: 6      Training Loss: 26136.892130      Validation Loss: 341
5.109824

Epoch: 7      Training Loss: 25605.508467      Validation Loss: 331
0.772789
    Saving model ...
Epoch: 8      Training Loss: 25193.762247      Validation Loss: 328
0.700325
    Saving model ...
Epoch: 9      Training Loss: 24664.896971      Validation Loss: 321
8.254964
    Saving model ...
Epoch: 10     Training Loss: 24299.343431      Validation Loss: 318
1.319473
    Saving model ...
```

```
Out[79]: <All keys matched successfully>
```

## Note2:

A bit better but start overfitting after epoche 9, I am going to change it as bellow. I test it for 50 epoch to see hoe it performs.

```

In [96]: import os
from torchvision import datasets

### TODO: Write data loaders for training, validation, and test sets
## Specify appropriate transforms, and batch_sizes
transform_resize = 224
transform_crop = 224
data_directory = "dogImages"

print("load image data ... ")
# define transforms for the training data and testing data

train_transforms = transforms.Compose([transforms.Resize(transform_re
size),
                                     transforms.RandomHorizontalFli
p(),
                                     #transforms.CenterCrop(transform_cro
rm_crop),
                                     transforms.RandomResizedCrop(t
ransform_resize, scale=(0.08,1), ratio=(1,1)),
                                     #transforms.RandomHorizontalFl
ip(),
                                     #transforms.RandomVerticalFlip
(),
                                     #transforms.RandomRotation(2
0),
                                     transforms.ToTensor(),
                                     transforms.Normalize([0.485,
0.456, 0.406],
                                                         [0.229,
0.224, 0.225]]))

test_transforms = transforms.Compose([transforms.Resize(transform_res
ize),
                                     transforms.CenterCrop(transform
_crop),
                                     transforms.ToTensor(),
                                     transforms.Normalize([0.485, 0.
456, 0.406],
                                                         [0.229, 0.
224, 0.225]]))

# pass transforms in here, then run the next cell to see how the tran
sforms look
train_data = datasets.ImageFolder( data_directory + '/train', transfo
rm=train_transforms )
test_data = datasets.ImageFolder( data_directory + '/test', transform
=test_transforms )
valid_data = datasets.ImageFolder( data_directory + '/valid', transfo
rm=test_transforms )

# ---- print out some data stats ----
print('  Number of train images: ', len(train_data))
print('  Number of test images: ', len(test_data))
print('  Number of valid images: ', len(valid_data))

```

```
# -----  
  
trainloader = torch.utils.data.DataLoader( train_data, batch_size=32,  
shuffle=True )  
testloader = torch.utils.data.DataLoader( test_data, batch_size=16 )  
validloader = torch.utils.data.DataLoader( valid_data, batch_size=16  
)  
  
# create dictionary for all loaders in one  
loaders_scratch = {}  
loaders_scratch['train'] = trainloader  
loaders_scratch['valid'] = validloader  
loaders_scratch['test'] = testloader  
n_classes = len(train_data.classes)  
  
load image data ...  
    Number of train images: 6680  
    Number of test images: 836  
    Number of valid images: 835
```

In [ ]:

```
In [95]: model_scratch = train(50, loaders_scratch, model_scratch, optimizer_s
cratch,
                                criterion_scratch, use_cuda, 'none_augment/mode
l_augmentation2_scratch.pt')

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('none_augment/model_augmenta
tion2_scratch.pt'))
```



```

Epoch: 1      Training Loss: 33289.013680      Validation Loss: 406
9.169052
  Saving model ...
Epoch: 2      Training Loss: 32003.008823      Validation Loss: 396
8.229279
  Saving model ...
Epoch: 3      Training Loss: 31424.128082      Validation Loss: 391
8.502841
  Saving model ...
Epoch: 4      Training Loss: 30893.075760      Validation Loss: 387
6.353407
  Saving model ...
Epoch: 5      Training Loss: 30494.032368      Validation Loss: 381
6.702356
  Saving model ...
Epoch: 6      Training Loss: 30068.997829      Validation Loss: 376
9.562776
  Saving model ...
Epoch: 7      Training Loss: 29692.055294      Validation Loss: 373
4.023375
  Saving model ...
Epoch: 8      Training Loss: 29333.469135      Validation Loss: 371
2.007081
  Saving model ...
Epoch: 9      Training Loss: 28963.161345      Validation Loss: 367
0.662122
  Saving model ...
Epoch: 10     Training Loss: 28482.317898      Validation Loss: 367
5.888136

Epoch: 11     Training Loss: 28217.756271      Validation Loss: 360
2.349226
  Saving model ...
Epoch: 12     Training Loss: 27997.603096      Validation Loss: 356
9.642857
  Saving model ...
Epoch: 13     Training Loss: 27623.184799      Validation Loss: 363
2.852955

Epoch: 14     Training Loss: 27279.235981      Validation Loss: 344
6.853124
  Saving model ...
Epoch: 15     Training Loss: 26986.070410      Validation Loss: 344
2.554840
  Saving model ...
Epoch: 16     Training Loss: 26823.070984      Validation Loss: 336
5.308640
  Saving model ...
Epoch: 17     Training Loss: 26293.918552      Validation Loss: 336
9.692473

Epoch: 18     Training Loss: 26070.885666      Validation Loss: 338
9.864799

Epoch: 19     Training Loss: 25697.397066      Validation Loss: 332
7.614166
  Saving model ...

```

Epoch: 20 2.726606	Training Loss: 25595.160543	Validation Loss: 334
Epoch: 21 8.532246	Training Loss: 25270.672350	Validation Loss: 325
Saving model ...		
Epoch: 22 8.496114	Training Loss: 25016.625162	Validation Loss: 324
Saving model ...		
Epoch: 23 6.919896	Training Loss: 24679.960295	Validation Loss: 318
Saving model ...		
Epoch: 24 5.723070	Training Loss: 24533.861128	Validation Loss: 314
Saving model ...		
Epoch: 25 2.367277	Training Loss: 24426.135239	Validation Loss: 316
Epoch: 26 9.893327	Training Loss: 24135.673895	Validation Loss: 321
Epoch: 27 7.609233	Training Loss: 23964.198538	Validation Loss: 309
Saving model ...		
Epoch: 28 4.265840	Training Loss: 23745.290079	Validation Loss: 302
Saving model ...		
Epoch: 29 8.837419	Training Loss: 23404.061167	Validation Loss: 309
Epoch: 30 9.051152	Training Loss: 23170.097200	Validation Loss: 312
Epoch: 31 9.955036	Training Loss: 23070.647598	Validation Loss: 326
Epoch: 32 8.775061	Training Loss: 22898.501955	Validation Loss: 317
Epoch: 33 7.862948	Training Loss: 22752.647522	Validation Loss: 300
Saving model ...		
Epoch: 34 2.860757	Training Loss: 22473.109957	Validation Loss: 292
Saving model ...		
Epoch: 35 1.648133	Training Loss: 22300.828423	Validation Loss: 313
Epoch: 36 2.139578	Training Loss: 22094.458605	Validation Loss: 290
Saving model ...		
Epoch: 37 2.403932	Training Loss: 21998.592911	Validation Loss: 291
Epoch: 38 2.105658	Training Loss: 21841.727213	Validation Loss: 290
Saving model ...		

```

Epoch: 39      Training Loss: 21736.855339   Validation Loss: 285
4.788197
  Saving model ...
Epoch: 40      Training Loss: 21429.024645   Validation Loss: 277
4.516261
  Saving model ...
Epoch: 41      Training Loss: 21196.958555   Validation Loss: 297
7.358753

Epoch: 42      Training Loss: 21341.049385   Validation Loss: 285
9.885207

Epoch: 43      Training Loss: 20857.280237   Validation Loss: 291
0.965622

Epoch: 44      Training Loss: 20829.968506   Validation Loss: 278
5.109583

Epoch: 45      Training Loss: 20580.846634   Validation Loss: 270
9.397722
  Saving model ...
Epoch: 46      Training Loss: 20467.874403   Validation Loss: 289
6.407328

Epoch: 47      Training Loss: 20421.746988   Validation Loss: 282
4.337608

Epoch: 48      Training Loss: 20304.645758   Validation Loss: 273
9.622649

Epoch: 49      Training Loss: 20211.557783   Validation Loss: 273
1.664424

Epoch: 50      Training Loss: 19991.709316   Validation Loss: 273
4.223965

```

**Out[95]:** <All keys matched successfully>

### Note3:

That is great! both are going down. I also observed that both start at lower error values now comparing previous trainings . It was a bit trickier than what I thought at firsts and time consuming too! next I am going to tweak the learning rate and find a semi optimal one.

### Tune CNN and hyper parameters:

I am going to train/test the model with a range of learning rates that seems reasonable to me... I am going to do it manually but I am sure there should be better ways like gridsearch.. I do it for just 10 epoches to get a sense , probably it is a good idea to do it for more epoches....

```
In [97]: param_learning_rate = 0.001 # I tweaked it manually
criterion_scratch = nn.CrossEntropyLoss() # Based on pytorch recommendation for multi class classification with high precision and recall

### TODO: select optimizer
optimizer_scratch = optim.SGD(model_scratch.parameters(), lr=param_learning_rate)
```

```
In [98]: model_scratch = train(10, loaders_scratch, model_scratch, optimizer_scratch,
                                criterion_scratch, use_cuda, 'learning_rate/model_scratch_001.pt')

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('learning_rate/model_scratch_001.pt'))
```

```
Epoch: 1      Training Loss: 19419.441353      Validation Loss: 257
1.991577
Saving model ...
Epoch: 2      Training Loss: 19154.762239      Validation Loss: 256
2.179347
Saving model ...
Epoch: 3      Training Loss: 19072.971766      Validation Loss: 255
3.935386
Saving model ...
Epoch: 4      Training Loss: 19076.901560      Validation Loss: 254
6.110691
Saving model ...
Epoch: 5      Training Loss: 19132.224756      Validation Loss: 255
5.562618

Epoch: 6      Training Loss: 18983.615808      Validation Loss: 258
1.055560

Epoch: 7      Training Loss: 18937.071451      Validation Loss: 254
2.892222
Saving model ...
Epoch: 8      Training Loss: 18976.418844      Validation Loss: 253
5.450424
Saving model ...
Epoch: 9      Training Loss: 19064.935646      Validation Loss: 256
3.680194

Epoch: 10     Training Loss: 18820.238550      Validation Loss: 255
7.209637
```

```
Out[98]: <All keys matched successfully>
```

It looks like now it learn faster, I am going to decrease learning rate even more.

```
In [ ]: param_learning_rate = 0.0005 # I tweaked it manually
model_scratch = train(10, loaders_scratch, model_scratch, optimizer_s
cratch,
                    criterion_scratch, use_cuda, 'learning_rate/mod
el_scratch_0005.pt')

# load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('learning_rate/model_scratch
_0005.pt'))
```

```
Epoch: 1      Training Loss: 18849.856522      Validation Loss: 255
8.956687
  Saving model ...
Epoch: 2      Training Loss: 18743.703451      Validation Loss: 254
6.099226
  Saving model ...
Epoch: 3      Training Loss: 18931.380140      Validation Loss: 255
0.586272

Epoch: 4      Training Loss: 18761.624598      Validation Loss: 255
5.475226

Epoch: 5      Training Loss: 18753.853256      Validation Loss: 253
1.223555
  Saving model ...
Epoch: 6      Training Loss: 18978.795713      Validation Loss: 253
1.353685

Epoch: 7      Training Loss: 18639.263281      Validation Loss: 254
4.865511

Epoch: 8      Training Loss: 18841.711714      Validation Loss: 256
2.131870

Epoch: 9      Training Loss: 18755.938637      Validation Loss: 252
3.797336
  Saving model ...
Epoch: 10     Training Loss: 18749.533445      Validation Loss: 253
1.421881
```

```
Out[ ]: <All keys matched successfully>
```

## Note 4 :

Slightly better but not huge improvement. I am going to stop looking for more optimum learning rate and use 0.0005 as lr.

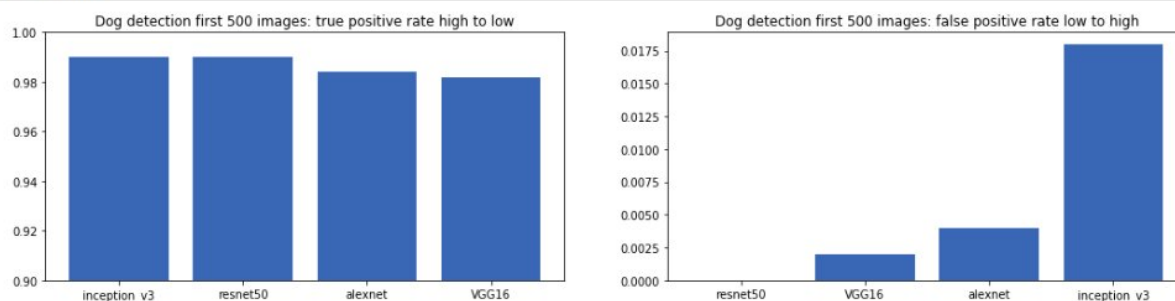
Adding more convolutional layer looks like a bit overkill, however replacing max global pooling layer with average global pooling could be considered.

## Transfer learning:

At the beginig of notebook. I have evaluated 4 models model\_names = ["alexnet","VGG16","resnet50","inception\_v3"] at turn out that resnet is doing a great job. I am going to use it as transfer learning model.All I need to do is to replace final linear layer.(model\_transfer.fc = nn.Linear(model\_transfer.fc.in\_features,n\_classes))

```
In [18]: idx_tpr = np.argsort(tpr)[::-1]
idx_fpr = np.argsort(fpr)

fig,ax = plt.subplots(1,2,figsize = (18,4))
ax[0].bar([model_names[i] for i in idx_tpr],tpr[idx_tpr])
ax[0].set_title("Dog detection first 500 images: true positive rate high to low")
ax[0].set_ylim((0.9,1))
ax[1].bar([model_names[i] for i in idx_fpr],fpr[idx_fpr])
ax[1].set_title("Dog detection first 500 images: false positive rate low to high")
plt.show()
```



**Answer:** Resnet50 outperforms other considering low fp and high tp.

### Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning yet!). and you must attain a test accuracy of at least 10%.

## Conclusion:

- As the dataset is small, it appears that data augmentation is needed to provide more training instance and avoid over fitting. Finding a good set/order of transformers was time consuming and tricky.
- Although resnet50 did the best job compare to others at my early test, I figured alexnet do actually detect the dogs better.
- I am sure there are better way to tune the hyper parameter as well as design the CNN classifier. I am still learning and hoping to get better understanding on deeper networks.

In [ ]: